

Media & Entertainment (OTT/Streaming) — Data Analyst/Scientist Domain Notes

01. Overview — What the Domain Is About & Industry Context

What is OTT/Streaming? OTT (Over-The-Top) refers to media/content delivered directly over the internet, bypassing traditional cable/satellite/broadcast distribution. This includes:

- **SVOD** (Subscription Video on Demand) — Netflix, Amazon Prime Video, Disney+, Hotstar
- **AVOD** (Advertising-based VOD) — YouTube, Pluto TV, Tubi, JioCinema (free tier)
- **TVOD** (Transactional VOD) — Apple TV rentals/purchases, Google Play Movies
- **Hybrid models** — Hotstar/JioCinema (ad-supported + subscription tiers), Netflix (ad-tier introduced)
- **Live streaming** — sports (IPL, NFL), news, live events
- **Audio streaming** (adjacent) — Spotify, Gaana, JioSaavn (similar analytics principles)

Industry Context (as of 2025-26):

- The industry has matured past pure subscriber-growth-at-any-cost; focus has shifted to **profitability, retention, and ARPU (Average Revenue Per User)** growth.
- Password-sharing crackdowns (Netflix), ad-tier launches (Netflix, Disney+, Amazon Prime), and bundling (e.g., telecom + OTT bundles common in India — Jio, Airtel) are major strategic levers.
- Content cost vs. content ROI is under heavy scrutiny — studios/platforms want to know which shows/movies actually drive subscriptions and retention, not just views.
- Regional/local content strategy is critical in markets like India, Latin America, Southeast Asia — global platforms increasingly invest in local-language originals.
- Live sports streaming (cricket, football) is a massive subscriber-acquisition and churn-prevention lever in markets like India (Hotstar/JioCinema during IPL/World Cup).
- Increasing use of AI for personalization, content recommendation, dynamic ad insertion, and content moderation.

Why Data Analysts/Scientists Matter Here: Streaming platforms generate massive behavioral data (every play, pause, skip, rewind, search, scroll) — the entire business (content investment, pricing, personalization, marketing spend) is data-driven. Analysts sit at the intersection of **content strategy, product, marketing, and finance**.

02. Role of a Data Analyst — Day-to-Day Responsibilities

A Data Analyst/Scientist in OTT typically sits within one of these functional pods:

- **Content Analytics** (which titles perform, content ROI)
- **Subscriber/Growth Analytics** (acquisition, churn, retention)
- **Product Analytics** (app/UX behavior, A/B testing of features)
- **Marketing Analytics** (campaign performance, CAC, attribution)
- **Ad Sales Analytics** (for AVOD — impression delivery, yield)
- **Recommendation/Personalization** (data science-heavy, ML models)

Typical day-to-day tasks:

1. **Dashboard monitoring** — Check daily/weekly metrics (DAU/MAU, churn rate, new signups, watch-time) via BI tools (Tableau/Looker/Power BI); flag anomalies (e.g., sudden drop in streams for a region).
2. **Ad-hoc analysis requests** — Content team asks "Did Season 2 of Show X retain Season 1 viewers?" — pull data, analyze, present findings.
3. **A/B test analysis** — Evaluate experiments (e.g., new homepage layout, new recommendation algorithm, price change) for statistical significance and business impact.
4. **Churn/retention modeling** — Build/maintain models predicting which users are likely to cancel; segment users for win-back campaigns.
5. **Content performance reporting** — Post-launch reports for new releases (viewership curve, completion rate, binge behavior, comparison vs. similar titles).
6. **Cohort analysis** — Track how different subscriber cohorts (signup month, acquisition channel, plan type) behave over time.
7. **SQL querying** — Pulling data from data warehouses (Snowflake/BigQuery/Redshift) daily for reporting and ad-hoc asks.

8. **Collaborating with data engineering** — Requesting new event tracking, validating data pipelines, fixing data quality issues.
 9. **Presenting to stakeholders** — Weekly/monthly business reviews (WBR/MBR) with content, marketing, and leadership teams.
 10. **Forecasting** — Subscriber growth forecasts, revenue forecasts, content budget allocation recommendations.
-

03. Key Metrics & KPIs

Subscriber/Growth Metrics

- **DAU/MAU** (Daily/Monthly Active Users) and **DAU/MAU ratio** (stickiness)
- **Net Subscriber Adds** = New signups – Cancellations
- **Churn Rate** (Voluntary vs. Involuntary — e.g., failed payment)
- **Retention Rate** (Day 1/7/30 retention, cohort-based)
- **ARPU** (Average Revenue Per User) / **ARPPU** (per paying user)
- **CAC** (Customer Acquisition Cost) and **LTV** (Lifetime Value); **LTV:CAC ratio**
- **Trial-to-Paid Conversion Rate**
- **Reactivation Rate** (win-back of churned users)

Content/Engagement Metrics

- **Watch Time / Total Streaming Hours**
- **Completion Rate** (% of a title watched to completion)
- **Time to First Play** (how quickly after signup a user starts watching)
- **Binge Rate** (episodes watched in one sitting)
- **Content Stickiness** = views on Title X / total platform views
- **Drop-off Points** (where in an episode/movie viewers abandon)
- **Repeat Viewership** (rewatch rate)
- **Search-to-Play Conversion** (did search lead to a play)
- **Recommendation CTR** (click-through rate on recommended titles)

Streaming Quality / Tech Metrics

- **Buffering Ratio / Rebuffer Rate**
- **Video Start Failure (VSF) Rate**
- **Bitrate/Video Quality delivered**
- **App Crash Rate**
- **Time to First Frame (startup latency)**

Monetization/Ads Metrics (AVOD)

- **Impressions Delivered / Fill Rate**
- **CPM (Cost per Mille)**
- **Ad Completion Rate (VCR - Video Completion Rate)**
- **Ad Load per Hour**

Financial/Business Metrics

- **Content Cost per Hour Watched** (ROI proxy for content spend)
- **Gross Margin per Title**
- **Revenue by Plan Tier** (Basic, Standard, Premium, Ad-tier)
- **Regional Revenue Contribution**

04. Common Datasets / Data Sources / Fields

1. User/Subscriber Data

- `user_id`, `signup_date`, `plan_type`, `region`, `device_type`, `payment_method`, `subscription_status`, `cancellation_date`, `cancellation_reason`

2. Viewing/Playback Events (largest dataset — clickstream-style)

- `user_id`, `content_id`, `session_id`, `play_start_timestamp`, `play_end_timestamp`, `watch_duration_seconds`, `device`, `bitrate`, `buffering_events`, `pause_events`, `seek_events`, `completion_flag`

3. Content Metadata

- content_id, title, genre, language, content_type (movie/series/episode), season_number, episode_number, release_date, runtime, cast/crew, content_cost, licensing_type (original/licensed), content_rating

4. Search & Discovery Data

- search_query, search_result_clicked, recommendation_slot_position, recommendation_source (algorithm/editorial), homepage_row_impressions

5. Marketing/Acquisition Data

- campaign_id, channel (social/search/referral/affiliate), acquisition_cost, attribution_source, promo_code_used

6. Billing/Transaction Data

- transaction_id, amount, currency, payment_status, refund_flag, invoice_date

7. Ad Server Data (AVOD-specific)

- impression_id, ad_id, advertiser, ad_break_position, viewability, completion_flag, revenue

8. Device/Technical Logs

- device_type, os_version, app_version, crash_logs, network_type, CDN_node, error_codes

Typical Data Volume Note: Playback event data is often massive (billions of rows/day for large platforms) — heavily reliant on distributed systems (Spark, Hadoop) and columnar warehouses (Snowflake, BigQuery, Redshift) with partitioning by date/region.

05. Tools & Technologies Commonly Used

Category	Tools
Languages	SQL (heavy daily use), Python (pandas, numpy, scikit-learn), R (less common now)
Data Warehousing	Snowflake, Google BigQuery, Amazon Redshift
Big Data Processing	Apache Spark, Hadoop, Databricks (for large-scale playback event processing)

Category	Tools
ETL/Pipelines	Airflow, dbt, Fivetran
BI/Visualization	Tableau, Looker, Power BI, Amazon QuickSight
Event Tracking/Product Analytics	Amplitude, Mixpanel, Segment (CDP), Google Analytics
Experimentation	Optimizely, in-house A/B testing frameworks, Statsig
ML/Recommendation	TensorFlow, PyTorch, Spark MLlib (for large-scale collaborative filtering)
Streaming Infra (context, not analyst-owned)	CDNs (Akamai, CloudFront), video encoding (AWS Elemental)
Version Control/Collaboration	Git, Jupyter Notebooks, Confluence, Jira
Cloud	AWS, GCP, Azure (platform-dependent)

06. Real-World Use Cases & Analytical Problems

1. Churn Prediction & Prevention

- Identify users likely to cancel in next 30 days using behavioral signals (declining watch time, payment failures, low engagement with new content).
- Trigger targeted retention offers (discount, curated content email) for high-risk segments.

2. Content Investment ROI Analysis

- Given a \$20M content budget, which genre/language/format delivers best "cost per hour watched" or best subscriber retention lift?
- Compare originals vs. licensed content ROI.

3. Personalization & Recommendation Effectiveness

- Measure impact of recommendation algorithm changes on watch time, session length, and CTR via A/B tests.

4. Password-Sharing / Account-Sharing Crackdown Analysis (Netflix-style)

- Detect shared-account patterns (multiple locations/devices), estimate revenue opportunity from converting to paid extra members.

5. Pricing & Plan Optimization

- Analyze price elasticity across regions/tiers; test impact of price hikes on churn vs. revenue.

6. Ad-tier Performance (AVOD)

- Optimize ad load (ads per hour) balancing revenue vs. viewer experience/churn.

7. Regional Content Strategy

- Determine which regional/local-language content drives incremental subscriber growth in specific markets (e.g., regional content performance in India, LatAm).

8. Live Event Streaming Capacity & Engagement

- For major live sports events, forecast concurrent viewership for CDN/infra planning; analyze engagement drop-offs during buffering.

9. Cohort Retention Curves

- Compare retention curves of cohorts acquired via different channels (organic vs. paid vs. bundle) to guide marketing spend allocation.

10. Trailer/Thumbnail A/B Testing

- Test different thumbnails/trailers for a title to optimize click-through and play-start rate.

07. ML / Statistical Techniques Relevant to the Domain

Technique	Application
Logistic Regression / Gradient Boosting (XGBoost, LightGBM)	Churn prediction, trial-to-paid conversion prediction
Survival Analysis (Cox Proportional Hazards, Kaplan-Meier)	Time-to-churn modeling, subscriber lifetime estimation
Collaborative Filtering (Matrix Factorization, ALS)	Recommendation systems (user-item interaction matrices)

Technique	Application
Content-Based Filtering / Embeddings	Recommending similar titles using metadata/content embeddings
Deep Learning (Neural Collaborative Filtering, Two-Tower models)	Modern large-scale recommendation systems
Clustering (K-Means, hierarchical)	User segmentation for marketing/personalization
A/B Testing & Causal Inference	Feature rollouts, pricing tests, recommendation algo changes; use of t-tests, chi-square, sequential testing
Time Series Forecasting (ARIMA, Prophet, LSTM)	Subscriber growth forecasting, watch-time trend forecasting, capacity planning for live events
Anomaly Detection	Detecting streaming quality issues, fraud/account-sharing patterns
Uplift Modeling	Measuring incremental impact of retention offers/marketing campaigns (who to target for max ROI)
NLP (Sentiment Analysis, Topic Modeling)	Analyzing reviews/social media sentiment about a title

Data Analyst/Scientist — Full Project Lifecycle (Step-by-Step)

This is the end-to-end process an analyst follows on a typical OTT project — e.g., **"Reduce churn by 10% for the Standard plan tier."**

Step 1: Problem Definition & Stakeholder Alignment

- Meet with product/content/marketing stakeholders to understand the business question.
- Clarify: What decision will this analysis inform? What's the timeline? Who's the audience for results?
- Translate vague ask ("users are leaving") into a measurable objective ("reduce 30-day voluntary churn rate for Standard-tier users by X%").

Step 2: Metric Definition & Success Criteria

- Define exact metric formulas (e.g., churn rate = cancellations in period / active subscribers at start of period).
- Agree on baseline, target, and time window with stakeholders.
- Identify guardrail metrics (e.g., don't reduce churn at the cost of tanking revenue via excessive discounting).

Step 3: Data Collection & Exploration

- Identify relevant data sources (subscriber DB, playback events, billing, marketing attribution).
- Write SQL queries to pull raw data from the warehouse (Snowflake/BigQuery/Redshift).
- Perform **Exploratory Data Analysis (EDA)**: check distributions, missing values, outliers, seasonality (e.g., churn spikes right after big sporting events end).

Step 4: Data Cleaning & Feature Engineering

- Handle missing/inconsistent data (e.g., null cancellation reasons, device mismatches).
- Engineer features: recency/frequency of watch sessions, average session length, genre diversity watched, days since last login, number of support tickets, payment failure history.
- Join disparate datasets (subscriber, engagement, billing) into an analysis-ready table.

Step 5: Segmentation / Cohort Analysis

- Segment users by acquisition channel, plan type, region, tenure, content affinity.
- Build cohort retention curves to see where drop-off concentrates (e.g., Month 2-3 for annual-sports-driven signups).

Step 6: Hypothesis Formulation

- Form hypotheses: "Users with <2 hrs watch time in first week churn 3x more" or "Users who don't engage with recommendations churn faster."
- Validate hypotheses statistically (correlation analysis, chi-square tests for categorical relationships).

Step 7: Modeling (if applicable)

- Build a churn prediction model (logistic regression / XGBoost) using engineered features.

- Split data into train/test (time-based split is important here to avoid leakage — train on past, test on future).
- Evaluate using precision/recall/AUC-ROC (churn is often imbalanced — prioritize recall or F1 depending on business cost of false negatives vs. positives).

Step 8: Insight Generation & Interpretation

- Use SHAP values / feature importance to explain what drives churn (interpretability matters for stakeholder trust).
- Translate model outputs into business insight: "Low engagement in first 14 days is the #1 churn driver; users who don't complete onboarding flow are 4x more likely to churn."

Step 9: Recommendation & Experiment Design

- Propose interventions: personalized onboarding nudges, targeted content recommendations for at-risk users, proactive retention offers.
- Design an A/B test to validate the intervention (control vs. treatment group), define sample size, duration, and success metric ahead of time (power analysis).

Step 10: A/B Test Execution & Monitoring

- Work with engineering to implement the test (feature flag / experiment framework like Statsig/Optimizely).
- Monitor test health during rollout (sample ratio mismatch, data quality checks, no unexpected side effects on other metrics).

Step 11: Results Analysis

- After test concludes, analyze statistical significance (t-test/chi-square) and practical significance (is the lift big enough to matter commercially).
- Check for heterogeneous treatment effects (did it work better for some segments than others).

Step 12: Reporting & Presentation

- Build a dashboard (Tableau/Looker) for ongoing monitoring of the metric.
- Present findings to stakeholders — clear narrative: problem → analysis → insight → recommendation → impact.
- Use visualizations: retention curves, funnel charts, cohort heatmaps.

Step 13: Decision & Rollout

- If test succeeds, work with product to roll out the winning treatment to 100% of relevant users.
- Document learnings in a shared knowledge base for future experiments.

Step 14: Post-Launch Monitoring

- Continue tracking the target metric post-rollout to confirm sustained impact (not just a novelty effect).
- Watch for metric drift or unintended consequences over subsequent months.

Step 15: Iteration

- Use learnings to formulate the next hypothesis/experiment (continuous improvement loop — this is rarely a one-and-done project).

08. Resources — Reference Links, Datasets, Courses, Papers

Public Datasets (for practice)

- **Netflix Prize Dataset** (classic recommendation systems dataset) — search "Netflix Prize dataset Kaggle"
- **MovieLens Dataset** (GroupLens Research) — user ratings, widely used for recommendation system practice
- **TMDB (The Movie Database) API** — content metadata for practice projects
- **Kaggle: "Netflix Movies and TV Shows"** dataset — content catalog analysis practice
- **Kaggle: OTT platform subscription/churn datasets** — search "OTT churn prediction Kaggle"

Courses/Learning

- Coursera: "Recommender Systems Specialization" (University of Minnesota)
- Udemy/Coursera: A/B Testing and Experimentation courses (e.g., Udacity's "A/B Testing" course)
- Google Analytics Academy (for product/behavioral analytics fundamentals)
- DataCamp: "Customer Analytics and A/B Testing in Python"

Industry Reports/Reading

- Ampere Analysis, Parrot Analytics (content performance benchmarking reports)
- Netflix Tech Blog (netflixtechblog.com) — excellent real-world case studies on recommendation systems, experimentation, personalization
- Reports from Ormax Media / Media Partners Asia (India/APAC OTT market specifically)

Papers

- "The Netflix Recommender System: Algorithms, Business Value, and Innovation" (ACM TMIS)
- "Deep Neural Networks for YouTube Recommendations" (Google, 2016) — foundational recommendation systems paper

Note: For a Logistics-domain analyst profile, OTT/Streaming is a good adjacent domain to know because both rely heavily on cohort/funnel analysis, churn/attrition modeling, and A/B testing — the analytical muscle transfers even though the business context differs.