

Domain Notes: Sports Analytics

01. Overview — What the Domain Is About & Industry Context

Sports Analytics is the application of data science, statistics, and machine learning to sports performance, strategy, business operations, fan engagement, and media. It has grown from a niche practice (popularized by "Moneyball" in baseball) into a core function across nearly every professional sport and league worldwide.

Why it matters right now:

- Every major league (NBA, NFL, EPL, cricket boards, F1, etc.) now runs player-tracking systems generating millions of data points per game.
- Sports betting/fantasy sports is a multi-billion-dollar industry that runs entirely on predictive models.
- Broadcasters and OTT platforms use analytics to personalize content and optimize viewer engagement.
- Team valuations, sponsorships, and player contracts increasingly hinge on data-driven performance projections.
- Athlete health/load management (injury prevention) has become a competitive differentiator, driven by wearable sensor data.

Two broad sub-domains within Sports Analytics:

1. **On-field/performance analytics** — player performance, tactics, injury prediction, talent scouting.
2. **Off-field/business analytics** — ticket pricing, fan engagement, sponsorship ROI, media rights valuation, fantasy sports/betting markets.

Who hires for this domain:

- Professional sports teams/franchises (NBA, NFL, EPL clubs, IPL franchises)
- Leagues & governing bodies (ICC, FIFA, NBA League Office, ATP/WTa)
- Sports media & broadcasters (ESPN, Star Sports, Disney)
- Betting & fantasy sports companies (Dream11, DraftKings, FanDuel, Bet365)
- Sports technology/wearables companies (Catapult, Stats Perform, Second Spectrum)

- Sports management/agency firms (player valuation, contract negotiation support)
- Sponsorship/marketing analytics consultancies

Core problem domains:

1. Player performance evaluation & talent scouting
 2. In-game strategy & tactical analysis
 3. Injury prediction & load management
 4. Fan engagement & audience analytics
 5. Betting odds & fantasy sports modeling
 6. Sponsorship & media valuation
 7. Ticket pricing & stadium operations
 8. Broadcast/viewership analytics
-

02. Role of the Data Analyst / Data Scientist in This Domain

Day-to-day responsibilities typically include:

- **Data collection & ingestion:** Pulling data from tracking systems (GPS vests, optical tracking cameras), official league APIs, and third-party providers (Opta, Stats Perform, Sportradar).
- **Cleaning & structuring event data:** Sports data is often event-based (every pass, shot, pitch) and needs careful structuring into analyzable formats (possession chains, play-by-play sequences).
- **Building performance metrics:** Creating advanced metrics beyond basic box-score stats (e.g., Expected Goals in football, Player Efficiency Rating in basketball).
- **Video + data synchronization:** Tagging tracking data to match video footage for coaches to review specific plays.
- **Predictive modeling:** Win probability models, player performance projections, injury risk scoring, fantasy point projections.
- **Opponent/scouting analysis:** Preparing pre-match reports on opposition tendencies (e.g., preferred formations, set-piece patterns) for coaching staff.

- **Dashboard/report creation:** Building dashboards for coaches, GMs, and analysts — often needs to be usable by non-technical staff under time pressure (e.g., halftime adjustments).
- **Draft/recruitment analytics:** Statistical models to evaluate amateur/college players or scout undervalued players in the transfer market.
- **Fan/business analytics:** Analyzing ticket sales, merchandise, social media engagement, and viewership data to guide marketing and pricing decisions.
- **Communicating to non-technical stakeholders:** Coaches and executives often aren't statisticians — translating complex models into simple, actionable insights is a critical skill in this domain.

Typical job titles: Sports Data Analyst, Performance Analyst, Basketball/Football Operations Analyst, Quantitative Analyst (Betting/Fantasy), Sports Scientist (data-focused), Recruitment/Scouting Analyst, Fan Insights Analyst, Sports Data Scientist.

03. Key Metrics & KPIs in This Domain

Category	Metric/KPI	Description
Football (Soccer)	Expected Goals (xG)	Probability a shot results in a goal based on shot location/type
Football (Soccer)	Expected Assists (xA)	Probability a pass leads to a goal
Basketball	Player Efficiency Rating (PER)	Composite per-minute productivity metric
Basketball	True Shooting % (TS%)	Shooting efficiency accounting for 2pt, 3pt, and free throws
Basketball	Plus/Minus, Net Rating	Point differential while a player/lineup is on court
Cricket	Strike Rate / Economy Rate	Batting scoring speed / bowling runs-conceded rate
Cricket	Win Probability (ball-by-ball)	Live probability of match outcome updated per delivery

Category	Metric/KPI	Description
American Football	QBR (Total Quarterback Rating)	Composite quarterback performance metric
General Performance	Player Load / Distance Covered	GPS-tracked physical exertion metric
General Performance	Injury Risk Score	Probability of injury based on load/biomechanical data
Team Strategy	Possession %, PPDA (Passes per Defensive Action)	Tactical style indicators (football)
Fantasy/Betting	Projected Fantasy Points	Expected statistical output per player per game
Fantasy/Betting	Implied Probability (odds)	Betting-market-derived win/outcome probability
Business/Fan	Net Promoter Score (NPS)	Fan satisfaction/loyalty metric
Business/Fan	Average Attendance / Ticket Yield	Stadium revenue efficiency metrics
Business/Fan	Social Media Engagement Rate	Fan interaction metric across platforms
Media	Viewership/Ratings, Time Spent Watching	Broadcast/streaming engagement metrics

04. Common Datasets / Data Sources Encountered

Tracking & Event Data (Providers):

- Opta / Stats Perform — event-level football, cricket, basketball data
- Second Spectrum / SportVU — optical player-tracking data (NBA)
- Catapult / STATSports — GPS wearable data for player load
- Hawk-Eye — ball-tracking data (tennis, cricket)
- Sportradar — live odds and event data feeds

League/Official APIs:

- NBA Stats API, NFL Next Gen Stats, ESPN API
- Cricsheet (open cricket ball-by-ball data)
- FBref / Understat (football statistics, xG data)
- ATP/WTB official stats feeds

Wearables/Biometric Data:

- Heart rate, GPS distance/speed, accelerometer data from player-worn devices
- Sleep and recovery data (from wearables like Whoop, Oura in some team setups)

Betting/Fantasy:

- Odds-comparison platforms (Odds API, Betfair Exchange data)
- Historical fantasy scoring data (Dream11, DraftKings, FanDuel platforms)

Business/Fan Data:

- Ticketing system data (internal CRM, StubHub/ticketing partner data)
- Social media APIs (Twitter/X, Instagram engagement data)
- Broadcast/streaming platform viewership data (internal to broadcasters)

Common data formats: JSON (event streams), CSV (box scores), XML (older feed formats), video files synced with timestamped tracking data, relational databases for historical stats warehousing.

05. Tools & Technologies Commonly Used

Category	Tools
Programming	Python (primary), R (statistical modeling, especially in academia/baseball analytics)
Data manipulation	Pandas, NumPy
Visualization	Matplotlib, Seaborn, Plotly, Tableau, mplsoccer/matplotlib-based pitch/court visualizations

Category	Tools
Specialized sports libraries	mpIsoccer, klopPy (football event data), nbastatR, hoopR (basketball), cricketstats/cricPy (cricket)
Databases	PostgreSQL/MySQL for historical stat warehouses, MongoDB for semi-structured event data
ML/Statistics	Scikit-learn, XGBoost/LightGBM, Bayesian modeling (PyMC/Stan) for probabilistic ratings
Video analysis	SportScore, Hudl, Catapult software (video-tagging synced with data)
Tracking/Spatial analysis	Custom pitch-control/spatial models, OpenCV for basic video-derived tracking
Cloud/Big data	AWS/GCP for storing and processing high-volume tracking data (millions of frames per game)
Reporting/BI	Power BI, Tableau, custom web dashboards (Streamlit/Dash) for coach-facing tools
Version control/collab	Git/GitHub, Jupyter Notebooks

06. Real-World Use Cases & Analytical Problems

1. **Expected Goals (xG) Modeling:** Building a model that assigns a goal-scoring probability to every shot based on distance, angle, body part, and defensive pressure — used to evaluate finishing quality and team attacking performance.
2. **Win Probability Models:** Live in-game models updating a team's probability of winning after every play/ball/possession, used for broadcast graphics and in-game strategy.
3. **Player Recruitment/Scouting Models:** Statistical player-similarity models to find undervalued players in lower leagues who match the profile of successful players (classic "Moneyball" application).
4. **Injury Risk Prediction:** Using GPS load data and historical injury records to flag players at elevated injury risk, informing rest/rotation decisions.
5. **Optimal Lineup/Rotation Analysis:** Analyzing plus/minus and on/off court data to determine best-performing player combinations.

6. **Fantasy Sports Point Projection:** Predicting player fantasy points for upcoming games using recent form, matchup difficulty, and historical performance.
 7. **Betting Odds Modeling:** Building probabilistic models for match outcomes to set or evaluate betting lines (quantitative analyst roles in betting firms).
 8. **Set-Piece/Tactical Analysis:** Breaking down opposition set-piece routines (corners, free kicks) using tracking data to prepare defensive strategies.
 9. **Draft Analytics:** Projecting how college/amateur players will perform at the professional level using historical draft outcome data.
 10. **Ticket Pricing Optimization:** Dynamic pricing models based on opponent quality, day of week, weather, and historical demand patterns.
 11. **Fan Churn Prediction:** Identifying season-ticket holders at risk of not renewing, based on engagement and attendance patterns.
 12. **Broadcast Content Personalization:** Recommending highlight clips or camera angles to individual viewers based on viewing behavior.
-

07. ML / Statistical Techniques Relevant to the Domain

Spatial & Event-Based Modeling:

- Logistic Regression / Gradient Boosting (XGBoost) — widely used for xG models (binary outcome: goal or no goal) due to interpretability requirements for coaching staff.
- Pitch-control models (physics/probability-based spatial models) — estimating which team controls space at any point in play.

Player/Team Rating Systems:

- Bayesian rating systems (Elo, Glicko, TrueSkill) — for ranking teams/players and updating ratings after each match.
- Plus/Minus regression (Adjusted Plus-Minus, RAPM in basketball) — isolating individual player impact from team performance using regularized regression.

Time-Series & Sequential Modeling:

- Markov Chain models — modeling possession/point sequences (e.g., baseball at-bat outcomes, tennis point sequences) as state transitions.

- LSTM/Recurrent models — modeling sequential play data (e.g., predicting next play type) using historical sequences.

Injury/Load Modeling:

- Survival analysis — time-to-injury modeling based on cumulative load.
- Random Forest/Gradient Boosting classifiers — injury risk classification using biomechanical and load features.

Computer Vision:

- CNNs / pose-estimation models — extracting player position and movement data directly from broadcast video (used when tracking-system data isn't available).
- Object detection (YOLO-based models) — automated ball/player tracking from video feeds.

Recommendation & Personalization (fan/business side):

- Collaborative filtering — content/highlight recommendation systems for streaming platforms.
- Churn prediction models (Logistic Regression, Gradient Boosting) — fan retention analytics.

Optimization:

- Linear/Integer programming — lineup optimization for fantasy sports (maximizing projected points under salary cap constraints), and roster construction under salary-cap rules.

08. Reference Resources — Links, Datasets, Courses, Papers

Datasets/Portals:

- Cricsheet (free ball-by-ball cricket data): <https://cricsheet.org>
- FBref (football statistics incl. xG): <https://fbref.com>
- Understat (football xG data): <https://understat.com>
- Basketball-Reference: <https://www.basketball-reference.com>
- Pro-Football-Reference: <https://www.pro-football-reference.com>

- Kaggle sports datasets (search "NBA", "football events", "cricket")
- StatsBomb Open Data (free football event data): <https://github.com/statsbomb/open-data>

Libraries/Tooling:

- mplsoccer (Python football visualization): <https://mplsoccer.readthedocs.io>
- kloppy (standardized sports event data parsing)
- nbastatR, hoopR (R packages for basketball data)

Courses:

- Coursera: "Sports Analytics" specialization
- Udemy/DataCamp: "Introduction to Sports Analytics with R/Python"
- Massey/Rasmussen sports analytics courses (many universities now offer dedicated programs, e.g., MIT Sloan Sports Analytics content)

Conferences/Communities:

- MIT Sloan Sports Analytics Conference (papers/talks archive — a major industry benchmark event)
- Kaggle sports competitions (NFL Big Data Bowl, March Madness prediction competitions)
- r/sportsanalytics, Twitter/X sports analytics community (many practitioners share open research)

Foundational reading:

- "Moneyball" by Michael Lewis — foundational case study in sports analytics adoption
- StatsBomb/Opta technical blogs on xG methodology

Full Project Lifecycle: What a Data Analyst/Scientist Actually Does, Step by Step

Phase 1: Problem Framing & Stakeholder Alignment

1. Meet with the requesting stakeholder (head coach, GM, recruitment director, or business/marketing team) to understand the actual question — e.g., "Should we sign this striker?" or "Which lineup should we start in the playoffs?"

2. Clarify what decision the analysis will inform and the timeline (in-season analysis often needs same-day turnaround; recruitment analysis may have weeks).
3. Identify what data is available/permissible (league data-sharing rules can restrict what tracking data is accessible, especially for opponent teams).

Phase 2: Data Acquisition

4. Pull data from internal systems (team's own tracking/wearable data) and external providers (Opta, Stats Perform, league APIs) via API or scheduled data feeds.
5. For scouting/opponent analysis, acquire publicly available or licensed third-party event data for teams/players outside your own organization.
6. Sync video footage timestamps with event/tracking data timestamps if the analysis needs coach-reviewable clips.

Phase 3: Data Cleaning & Preprocessing

7. Standardize player/team IDs across multiple data providers (a major real-world pain point — different vendors use different naming/ID conventions).
8. Handle missing or misclassified events (e.g., tracking system losing a player off-screen, mislabeled event types).
9. Normalize data across seasons/rule changes (e.g., 3-point line distance changes, shot clock rule changes) so historical comparisons remain valid.
10. Structure raw event streams into analyzable units (possessions, plays, at-bats, rallies) with contextual features attached (score, time remaining, game situation).

Phase 4: Exploratory Data Analysis (EDA)

11. Visualize player/team performance trends over the season (form curves, rolling averages).
12. Create spatial visualizations (shot maps, heatmaps, pass networks) to understand tactical patterns.
13. Compare against league-average baselines/percentiles to contextualize whether a metric is actually good or bad.
14. Identify small-sample-size issues — a common trap in sports analytics is over-reacting to metrics based on very few games/events.

Phase 5: Feature Engineering

15. Build advanced contextual features: game state (score differential, time remaining), opponent strength adjustment, home/away effects, rest days between games.
16. Create rolling-window features (last 5/10 game form) and career-stage features (age curves, experience).
17. Engineer matchup-specific features for scouting (e.g., performance against left-handed vs right-handed pitchers, performance against specific defensive schemes).

Phase 6: Modeling

18. Choose the right technique for the problem (see Section 07): classification for injury risk/xG, regression for point projections, rating systems for player value, optimization for lineup/roster construction.
19. Split data respecting temporal order (train on past seasons, validate on most recent season — never randomly shuffle across time in sports data, since player performance evolves).
20. Build a simple baseline first (e.g., "player will perform at their season average") before complex models — sports metrics often have highly predictable baselines.
21. Train and tune models, incorporating known sports domain adjustments (e.g., regression to the mean for small-sample performances, strength-of-schedule adjustments).
22. Validate model outputs against domain expert intuition (coaches/scouts) — a model that contradicts well-known tactical wisdom needs re-examination.

Phase 7: Evaluation & Interpretation

23. Evaluate using appropriate metrics (log-loss/Brier score for win-probability models, RMSE for point projections, AUC for injury-risk classifiers).
24. Backtest models against historical outcomes (e.g., "would this recruitment model have correctly identified last year's breakout players?").
25. Quantify uncertainty and communicate confidence levels — especially important for recruitment decisions involving significant financial investment (transfer fees, contracts).

Phase 8: Reporting & Communication

26. Build visual, coach-friendly reports/dashboards — often simplified single-page visuals rather than complex statistical tables, since coaching staff need fast, actionable insight, sometimes at halftime.

27. Prepare scouting reports combining data insights with video clips for recruitment/coaching decisions.
28. Present findings to non-technical stakeholders using sports-native language and visuals (shot maps, radar/spider charts) rather than raw statistical output.

Phase 9: Deployment & In-Season Operationalization

29. Automate data pipelines so dashboards update automatically after each game/match (critical during the season when turnaround time matters).
30. Integrate live models (win probability, in-game fantasy projections) into broadcast graphics or team-facing live dashboards where relevant.
31. Set up alerts for significant deviations (e.g., a player's load exceeding a safe threshold, triggering an injury-risk flag to medical staff).

Phase 10: Continuous Improvement & Governance

32. Retrain/recalibrate models each season as rules, rosters, and playing styles evolve (models trained on old rule sets/tactics can degrade quickly).
 33. Track model performance against real outcomes over time (e.g., did recruited players actually perform as projected) and feed this back into future model iterations.
 34. Document methodology clearly, since sports analytics conclusions are often scrutinized by coaches, media, and ownership — transparency and reproducibility build trust in the analytics function.
 35. Stay current with evolving tracking technology (broadcast-video-derived tracking, new wearable sensors) since data quality/availability in this domain improves rapidly year over year.
-